

Data Processing in the Tidyverse

Using Administrative Data for
Clinical and Health Services Research

Overview

- Tidyverse plug
- Selecting columns with `select`
- Subsetting rows with `filter`
- Summarizing by group with `group_by` / `summarise`



R packages for data science

The tidyverse is an opinionated collection of **packages** designed for data science that share an underlying design philosophy, coding conventions, and data structures.

Install the complete tidyverse with

```
install.packages("tidyverse")
```



Overview

dplyr is a grammar of data manipulation, providing a consistent set of verbs that help you solve the most common data manipulation challenges:

- `mutate()` adds new variables that are functions of existing variables
- `select()` picks variables based on their names.
- `filter()` picks cases based on their values.
- `summarise()` reduces multiple values down to a single summary.
- `arrange()` changes the ordering of the rows.

These all combine naturally with `group_by()` which allows you to perform any operation “by group”. You can learn more

Links

Download

<https://cran.r-project.org/web/packages/dplyr/index.html>

Browse source

<https://github.com/tidyverse/dplyr>

Report a problem

<https://github.com/tidyverse/dplyr/issues>

Learn more



Layers

Geoms

A layer combines data, aesthetic mapping, a geom (geometric object), a stat (statistical transformation), and a position adjustment. Typically, you will create layers using a `geom_` function, overriding the default position and stat if needed.



`geom_abline()` `geom_hline()` `geom_vline()` Reference lines: horizontal, vertical, and diagonal



`geom_bar()` `geom_col()` `stat_count()` Bar charts



`geom_bin_2d()` `stat_bin_2d()` Heatmap of 2d bin counts



`geom_blank()` Draw nothing



`geom_boxplot()` `stat_boxplot()` A box and whiskers plot (in the style of Tukey)

Contents

Plot basics

Layers

Aesthetics

Scales

Guides: axes and legends

Facetting

Coordinate systems

Themes

Programming with ggplot2

Extending ggplot2

Vector helpers

Data

Autoplot and fortify

Selecting columns

- Specific columns in a data set can be chosen using `dplyr::select`
 - As with `mutate`, first argument is a data set, subsequent arguments are expressions, and the result is a new data set
- Offers a variety of ways to choose multiple columns at once, making it a significant improvement on approaches available in base R.

Selecting columns

By name

```
# Create a data set of discharge- and patient-specific identifiers.  
disch_identifiers <- select(corelp, KEY, VisitLink, DaysToEvent, LOS)  
colnames(disch_identifiers)
```

```
## [1] "KEY"          "VisitLink"    "DaysToEvent" "LOS"
```

Selecting columns

By a character vector

```
# Create a data set containing the variables related to patient  
# residence.  
vars_patient_residence <- c(  
  "PSTATE", "PSTCO", "PSTCO2", "PSTCO_GEO", "ZIP")  
disch_patient_residence <- select(  
  corelp, all_of(vars_patient_residence))  
colnames(disch_patient_residence)
```

```
## [1] "PSTATE"    "PSTCO"     "PSTCO2"    "PSTCO_GEO" "ZIP"
```


Selecting columns

By a common prefix

```
# Create a data set consisting of the procedure code columns.  
disch_prccodes <- select(corelp, starts_with("I10_PR"))  
colnames(disch_prccodes)
```

```
## [1] "I10_PR1" "I10_PR2" "I10_PR3" "I10_PR4" "I10_PR5" "I10_PR6"  
## [7] "I10_PR7" "I10_PR8" "I10_PR9" "I10_PR10" "I10_PR11" "I10_PR12"  
## [13] "I10_PR13" "I10_PR14" "I10_PR15" "I10_PR16" "I10_PR17" "I10_PR18"  
## [19] "I10_PR19" "I10_PR20" "I10_PR21" "I10_PR22" "I10_PR23" "I10_PR24"  
## [25] "I10_PR25" "I10_PR26" "I10_PR27" "I10_PR28" "I10_PR29" "I10_PR30"  
## [31] "I10_PR31"
```

Selecting columns

By a common suffix

```
# Select all of the "from source" columns (i.e., columns whose values  
# were not edited by HCUP).  
disch_from_source <- select(corelp, ends_with("_X"))  
colnames(disch_from_source)
```

```
## [1] "DISP_X"    "LOS_X"     "PAY1_X"    "RACE_X"    "TOTCHG_X"
```

Selecting columns

By a common sequence of characters somewhere in the names

```
# Select all columns with income-related information.  
disch_income <- select(corelp, contains("INC"))  
colnames(disch_income)
```

```
## [1] "MEDINCSTQ" "ZIPINC_QRTL"
```

Selecting columns

By sequence based on column order

- Inspired by shorthand for sequence of integers: $1:3 \equiv c(1,2,3)$

```
# Select `VisitLink` and `DSHOSPID` and everything in between.
disch_VisitLink_to_DSHOSPID <- select(corelp, VisitLink:DSHOSPID)
colnames(disch_VisitLink_to_DSHOSPID)
```

```
## [1] "VisitLink" "AGE" "AHOUR" "ATYPE"
## [5] "AWEEKEND" "DaysToEvent" "DIED" "DISP_X"
## [9] "DISPUB04" "DISPUNIFORM" "DQTR" "DSHOSPID"
```

- Using column positions can be risky
- Suppose we want to select all diagnosis code columns

```
disch_all_dx_codes_wrong <- select(
  core1p, I10_DX_Admitting:I10_DX34, I10_ECAUSE1:I10_ECAUSE6)
colnames(disch_all_dx_codes_wrong)
```

```
##      [1] "I10_DX_Admitting"      "I10_DX1"              "I10_DX2"
##      [6] "I10_DX5"              "I10_DX6"              "I10_DX7"
##     [11] "I10_DX10"             "I10_DX11"             "I10_DX12"
##     [16] "I10_DX15"             "I10_DX16"             "I10_DX17"
##     [21] "I10_DX20"             "I10_DX21"             "I10_DX22"
##     [26] "I10_DX25"             "I10_DX26"             "I10_DX27"
##     [31] "I10_DX30"             "I10_DX31"             "I10_ECAUSE1"
##     [36] "I10_ECAUSE4"          "I10_ECAUSE5"          "I10_ECAUSE6"
##     [41] "I10_NPR"              "I10_PR1"              "I10_PR2"
##     [46] "I10_PR5"              "I10_PR6"              "I10_PR7"
##     [51] "I10_PR10"             "I10_PR11"             "I10_PR12"
##     [56] "I10_PR15"             "I10_PR16"             "I10_PR17"
##     [61] "I10_PR20"             "I10_PR21"             "I10_PR22"
##     [66] "I10_PR25"             "I10_PR26"             "I10_PR27"
##     [71] "I10_PR30"             "I10_PR31"             "KEY"
##     [76] "MEDINCSTQ"            "PAY1"                  "PAY1_X"
##     [81] "PL_RUCC"              "PL_UIC"               "PL_UR_CAT4"
##     [86] "POA_Hosp_Edit1"       "POA_Hosp_Edit3_Value" "PointOfOriginUB04"
##     [91] "PRDAY3"               "PRDAY4"               "PRDAY5"
##     [96] "PRDAY8"               "PRDAY9"               "PRDAY10"
```

Selecting columns

By sequence based on column order

- Different years of data have different numbers of diagnosis codes
- 1% sample data set was formed by stacking the annual CORE files
 - I10_DX_Admitting-I10_DX31 contiguous but not with I10_DX32-I10_DX34
- In general, selection based on column names is more reliable
- Regardless, always check your selection with `str`, `colnames`, or by printing or Viewing the resulting data set

Selecting columns

By sequence based on column order

```
disch_all_dx_codes_right <- select(  
  corelp, starts_with("I10_DX"), starts_with("I10_ECAUSE"))  
colnames(disch_all_dx_codes_right)
```

```
## [1] "I10_DX_Admitting" "I10_DX1"          "I10_DX2"  
## [4] "I10_DX3"          "I10_DX4"          "I10_DX5"  
## [7] "I10_DX6"          "I10_DX7"          "I10_DX8"  
## [10] "I10_DX9"          "I10_DX10"         "I10_DX11"  
## [13] "I10_DX12"         "I10_DX13"         "I10_DX14"  
## [16] "I10_DX15"         "I10_DX16"         "I10_DX17"  
## [19] "I10_DX18"         "I10_DX19"         "I10_DX20"  
## [22] "I10_DX21"         "I10_DX22"         "I10_DX23"  
## [25] "I10_DX24"         "I10_DX25"         "I10_DX26"  
## [28] "I10_DX27"         "I10_DX28"         "I10_DX29"  
## [31] "I10_DX30"         "I10_DX31"         "I10_DX32"  
## [34] "I10_DX33"         "I10_DX34"         "I10_ECAUSE1"  
## [37] "I10_ECAUSE2"      "I10_ECAUSE3"      "I10_ECAUSE4"  
## [40] "I10_ECAUSE5"      "I10_ECAUSE6"
```

Selecting columns

By negation

- i.e., "select everything except..."

```
# Create a data set excluding all columns related to diagnosis and
# procedures codes, including POA codes and the `PRDAY`n fields.
disch_no_dx_or_pr_codes <- select(
  corelp,
  !c(
    starts_with("I10_"), starts_with("DXPOA"), starts_with("E_POA"),
    starts_with("PRDAY"))
)
colnames(disch_no_dx_or_pr_codes)
```

```
## [1] "VisitLink"      "AGE"
## [3] "AHOUR"          "ATYPE"
## [5] "AWEEKEND"       "DaysToEvent"
## [7] "DIED"           "DISP_X"
## [9] "DISPUB04"       "DISPUNIFORM"
## [11] "DQTR"           "DSHOSPID"
```


Subsetting rows

- Aside from the column exclusion, `!` is often used in three cases:

1. Select nonmissing values

- `!is.na(x)`

2. Checking for non-equality

- `x != 0`

3. Select values *not in* a specific set of values

- `!(x %in% y)`
- Can be difficult to read if part of complex expression, so may want to use special 'not in' operator `x %nin% y`; this requires loading Hmisc with `library(Hmisc)`

Subsetting rows

- Relies on expressions that return TRUE or FALSE
- Used several of these types of expressions when defining variables in previous lecture:
 - `is.na(LOS)`
 - `AGE >= 18`
 - `PAY1 == 1`
 - `TRAN_IN %in% c(1, 2)`

Subsetting rows

- Tidyverse subsetting uses `filter`
- Can combine multiple expressions or write them as separate arguments
 - The arguments are implicitly combined using `&`

```
disch_adults1 <- filter(core1p, !is.na(AGE) & AGE >= 18)  
disch_adults2 <- filter(core1p, !is.na(AGE), AGE >= 18)
```

- Expressions connected by 'or' have to be combined

```
disch_transfer_in_or_out <- filter(  
  core1p, TRAN_IN %in% c(1, 2) | TRAN_OUT %in% c(1, 2))
```

Grouping data

- We can tell dplyr we want to process data by a grouping factor using `group_by`
 - Doesn't directly change the data set, just how other functions work

Summarizing by group

- `group_by` is most often used with `summarise` to compute summary statistics within each group

```
# Create a data set containing the number of discharges for each  
# patient.  
disch_grouped_by_patient <- group_by(corelp, VisitLink)  
p_discharge_counts <- summarise(  
  disch_grouped_by_patient, discharge_count = n())
```

- `group_by` paired with `summarise` changes the level of organization of the data—here from discharge level to patient level

Summarizing by group

```
disch_grouped_by_patient
```

```
## # A tibble: 102,733 × 194
## # Groups:   VisitLink [48,085]
##   VisitLink  AGE AHOUR ATYPE AWEKEND DaysToEvent  DIED DISP_X
##   <dbl> <dbl> <dbl> <dbl>    <dbl>    <dbl> <dbl> <chr>
## 1 15965139   74  1400     1         0    18739     0 01
## 2 16050018   77  1300     1         0    18655     0 03
## 3 16050018   77  2300     1         1    18623     0 03
## 4  2092107   32  2300     1         0    17883     0 01
## 5  8726364   66  1600     1         0    17252     0 06
## 6  390742    61  2300     1         0    18869     0 06
## 7  390742    61     0     1         1    18844     0 01
## 8 24417038   62  2200     1         1    19810     0 01
## 9 24344014   67   600     1         0    20164     0 01
## 10 22892274    9   900     3         0    16852     0 01
## # ... with 102,723 more rows, and 186 more variables: DIS PUB04 <dbl>,
## #   DISP UNIFORM <dbl>, DQTR <dbl>, DSHOSP ID <chr>, DXPOA1 <chr>,
## #   DXPOA2 <chr>, DXPOA3 <chr>, DXPOA4 <chr>, DXPOA5 <chr>,
## #   DXPOA6 <chr>, DXPOA7 <chr>, DXPOA8 <chr>, DXPOA9 <chr>,
## #   DXPOA10 <chr>, DXPOA11 <chr>, DXPOA12 <chr>, DXPOA13 <chr>,
## #   DXPOA14 <chr>, DXPOA15 <chr>, DXPOA16 <chr>, DXPOA17 <chr>,
```

Summarizing by group

```
p_discharge_counts
```

```
## # A tibble: 48,085 × 2
##   VisitLink discharge_count
##   <dbl>          <int>
## 1         365            3
## 2         626            3
## 3         641            4
## 4         769            1
## 5         846            8
## 6         951            9
## 7        1304            1
## 8        1411            2
## 9        1463            3
## 10       2556            1
## # ... with 48,075 more rows
```

The pipe `%>%` operator

- Many packages in the Tidyverse are designed to work with the pipe operator, `%>%`
- Takes the result of the expression on the 'left-hand side' and passes it as the first argument to the expression on the 'right-hand side'.
- Useful because each of the tidyverse functions we've looked at so far—`mutate`, `select`, `filter`, `group_by`, and `summarise`—all take a data set as their first argument and produce a new data set as output

The pipe %>% operator

So we can write

```
disch_identifiers <- select(corelp, KEY, VisitLink, DaysToEvent, LOS)
```

as

```
disch_identifiers <- corelp %>%  
  select(KEY, VisitLink, DaysToEvent, LOS)
```

- Most useful when needing to 'chain' several commands together but don't want to save intermediate results or write a nested set of function calls

Suppose we want the procedure code and procedure day columns for all adults starting in 2016. We could write

```
disch_adults <- filter(core1p, !is.na(AGE), AGE >= 18, YEAR >= 2016)
disch_adults_pra <- select(
  disch_adults,
  KEY, VisitLink, DaysToEvent, LOS, starts_with("I0_PR"),
  starts_with("PRDAY"))
```

or

```
disch_adults_prb <-
  select(
    filter(
      core1p,
      !is.na(AGE), AGE >= 18, YEAR >= 2016
    ),
    KEY, VisitLink, DaysToEvent, LOS, starts_with("I0_PR"),
    starts_with("PRDAY")
  )
```

The pipe %>% operator

But a bit clearer would be

```
disch_adults_prc <- corelp %>%  
  filter(!is.na(AGE), AGE >= 18, YEAR >= 2016) %>%  
  select(  
    KEY, VisitLink, DaysToEvent, LOS, starts_with("I0_PR"),  
    starts_with("PRDAY")  
  )
```